

A hand holding a Rubik's cube, with the text overlaid on the image.

# CS 957, System-2 AI Neuro-Symbolic AI

Mohammad Hossein Rohban | Feb 2025

Sharif University of Technology

---

# Interpretable Neural-Symbolic Concept Reasoning

---

**Pietro Barbiero<sup>\*1</sup> Gabriele Ciravegna<sup>\*2</sup> Francesco Giannini<sup>\*3</sup> Mateo Espinosa Zarlenga<sup>1</sup>**  
**Lucie Charlotte Magister<sup>1</sup> Alberto Tonda<sup>4</sup> Pietro Lió<sup>1</sup> Frederic Precioso<sup>2</sup> Mateja Jamnik<sup>1</sup>**  
**Giuseppe Marra<sup>\*5</sup>**


# Motivation

Given an input to be classified.

Deep models are **accurate** but are **blackbox** and are not interpretable.

Inherently interpretable models are **interpretable** but are **not as accurate**.

Can we have **both**?



Courtesy: Author's slides

# Motivation (cont.)

Want to build an


- inherently interpretable
- accurate
- robust

model.

# Concepts

Each input can be decomposed into a set of **concepts** that **together** **characterize** the input semantic.

apple = red + round



Courtesy: Author's slides

# Initial Idea

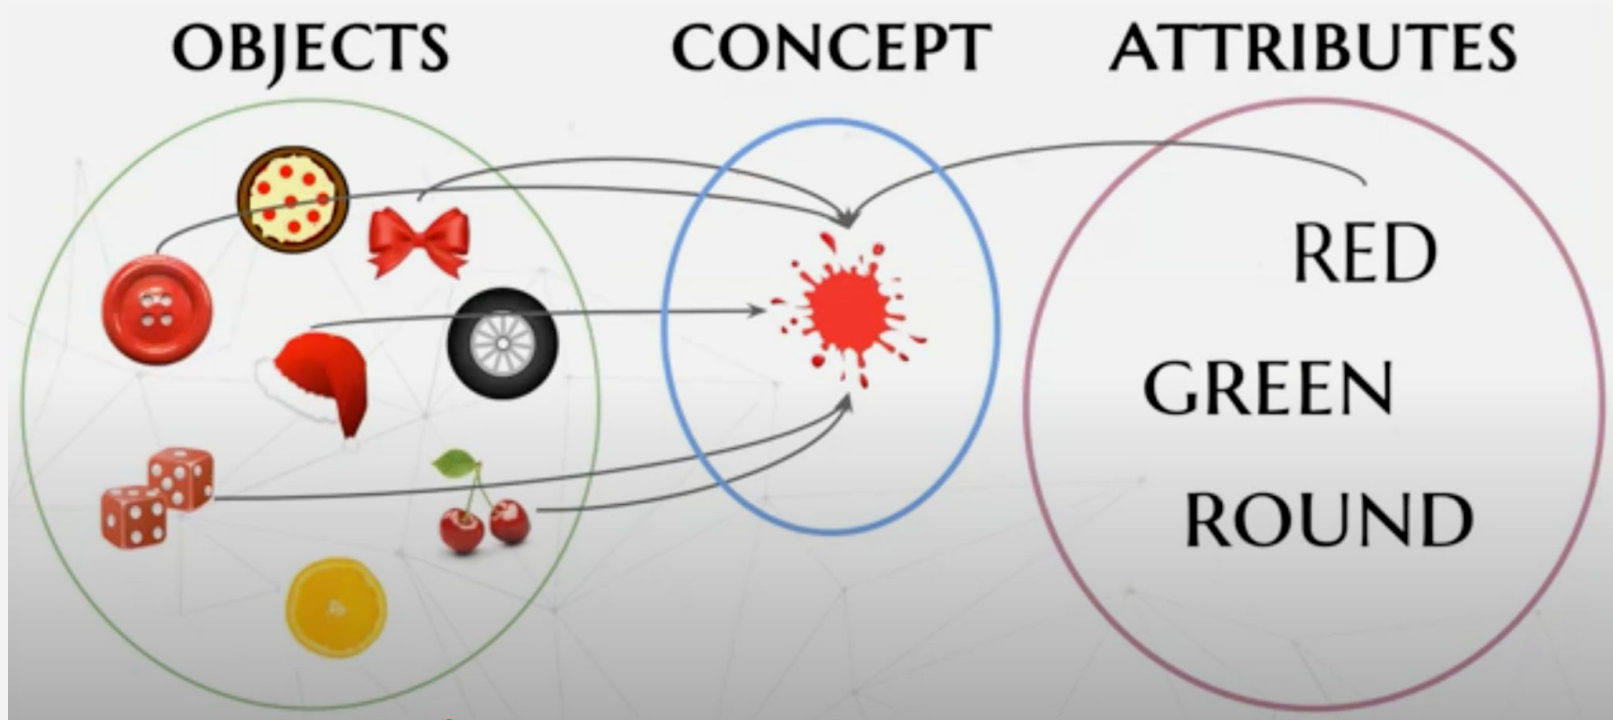
Describe each class as logical AND of concepts:

Apple = Red  $\wedge$  Round

This is definitely interpretable, but a whole lot of questions arise:

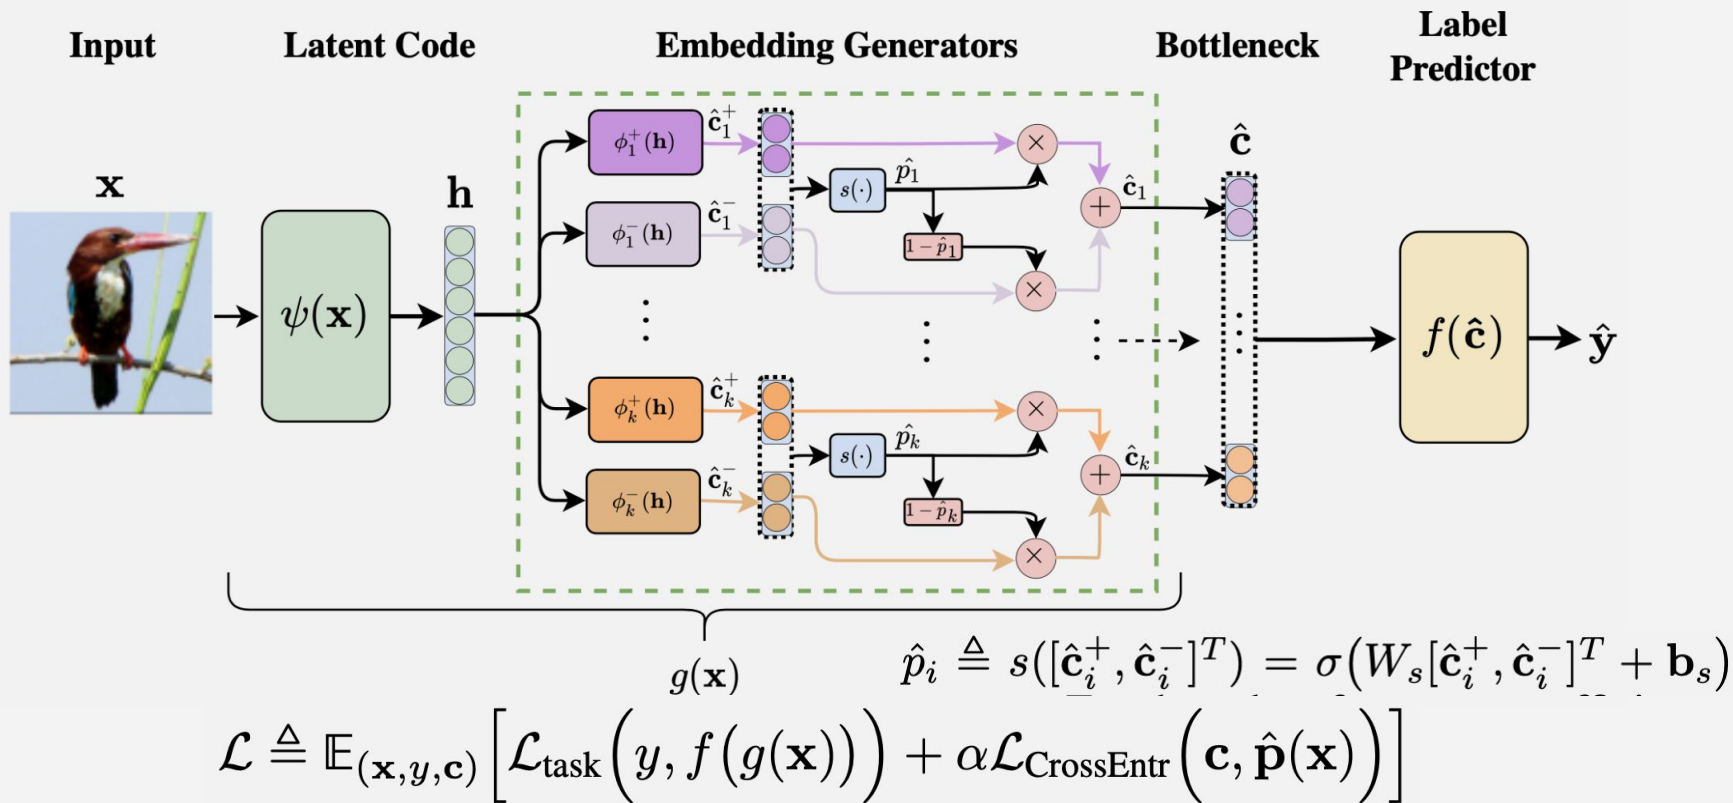
- Where do **concepts** come from?
- Where do **logical rules** come from?
- **Crisp** rules tend to result in **inaccurate** classification.

But what are concepts?



Courtesy: Author's slides

# We can try to learn them ...





# We we have up to so far?

For each concept  $i$ , an embedding vector  $\mathbf{c}_i$ .

A **deep model** that takes as input  $x$ , and outputs **existence**, in form of confidence, of each concept  $i$ .

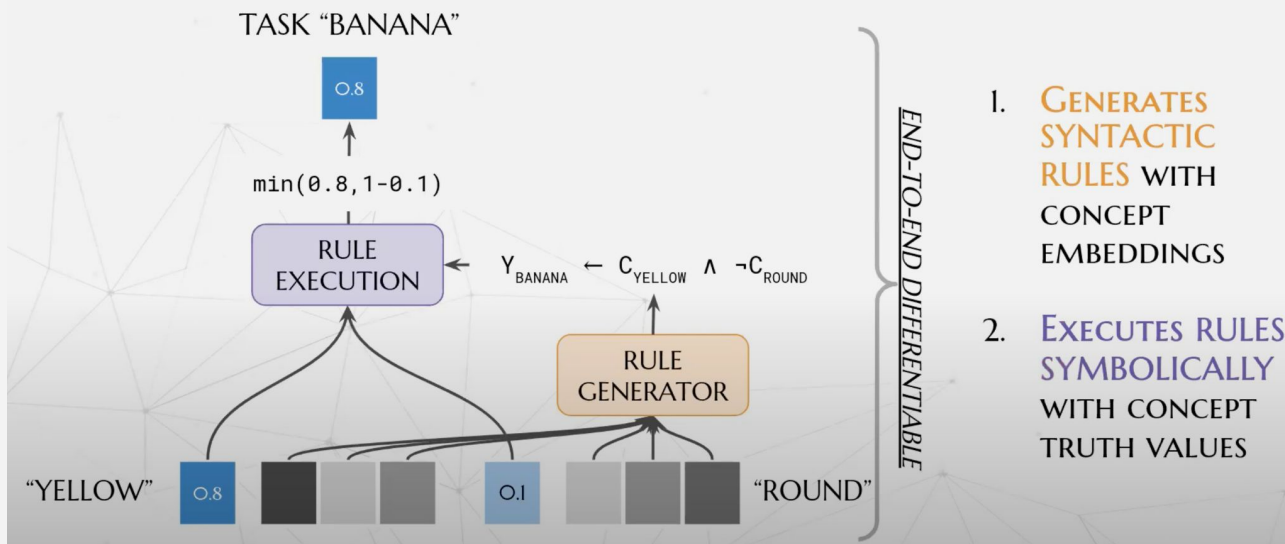
Let's call it  $g(x)$ .

How to form **rules** then?

# Rule Learning

Use a **neural network** to output a rule.

Then, rules are **executed**.



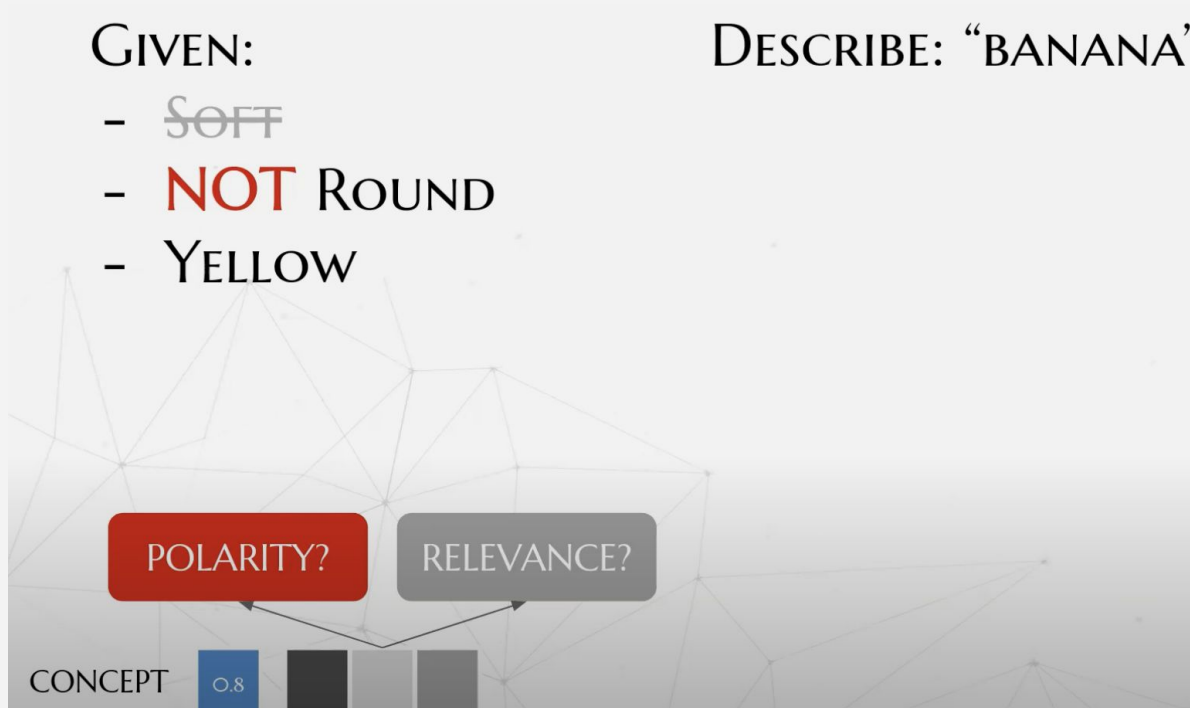
Courtesy: Author's slides

# Rule Learning (cont.)

GIVEN:

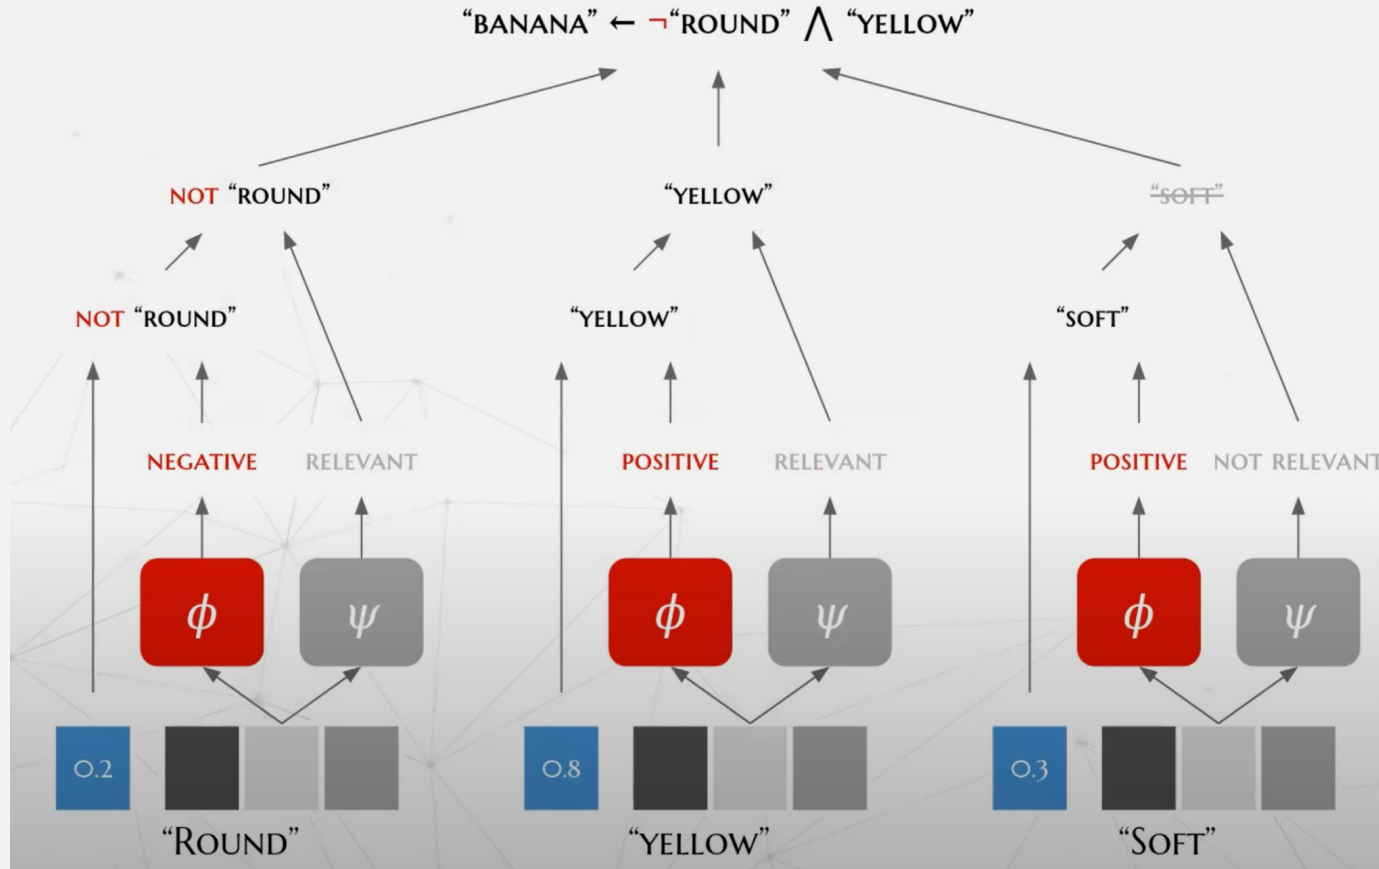
- ~~SOFT~~
- **NOT** ROUND
- YELLOW

DESCRIBE: “BANANA”



Courtesy: Author's slides

# Rule Learning (cont.)



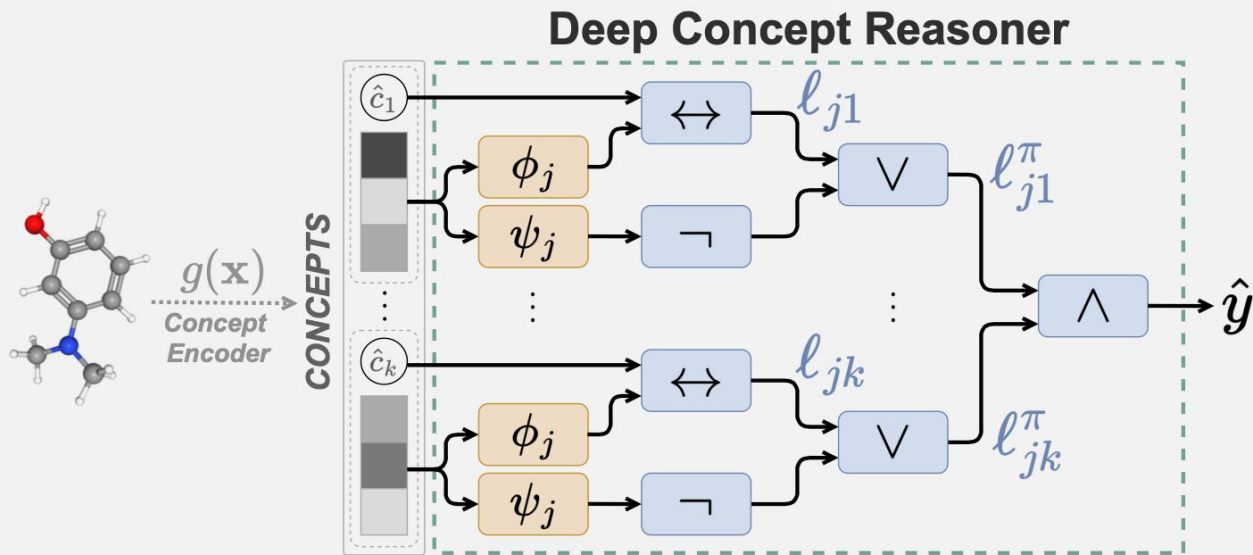
Courtesy: Author's slides

# Rule Learning (cont.)

Literal  $l_{ij}$ : for a given  $x$ , concept  $i$  status is in favor of class  $j$ :  $l_{ji} = (\phi_j(\hat{\mathbf{c}}_i) \Leftrightarrow \hat{\mathbf{c}}_i)$

Literal  $l_{ij}^r$ : if concept  $i$  is relevant to class  $j$ , then  $l_{ij}$ , otherwise 1.

$$l_{ji}^r = (\psi_j(\hat{\mathbf{c}}_i) \Rightarrow l_{ji}) = (\neg\psi_j(\hat{\mathbf{c}}_i) \vee l_{ji})$$

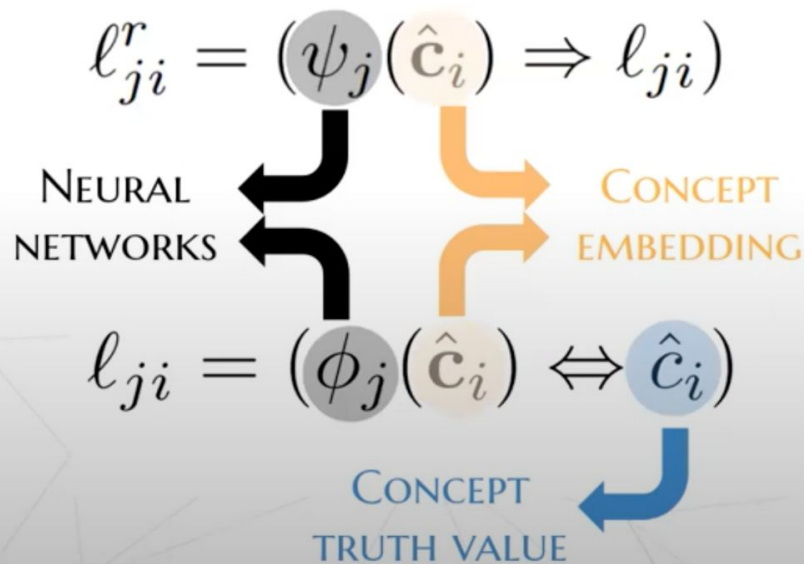


PREDICTION

RELEVANCE

POLARITY

$$\hat{y}_j = \bigwedge_{i=1}^k \ell_{ji}^r$$



# Making it differentiable

Make discontinuous logical binary operators continuous.

Fuzzy logic: each statement has a degree of truth

Logical operators on statements are functions of their truth values.

op \ t-norm	Product	Lukasiewicz	Gödel
$x \wedge y$	$x \cdot y$	$\max(0, x + y - 1)$	$\min(x, y)$
$x \vee y$	$x + y - x \cdot y$	$\min(1, x + y)$	$\max(x, y)$
$\neg x$	$1 - x$	$1 - x$	$1 - x$
$x \Rightarrow y$	$x \leq y ? 1 : \frac{y}{x}$	$\min(1, 1 - x + y)$	$x \leq y ? 1 : y$

# Differentiable Expression

Adopting Godel's t-norm:

$$\begin{aligned}\ell_{ji} &= \phi_j(\hat{\mathbf{c}}_i) \Leftrightarrow \hat{c}_i = (\phi_j(\hat{\mathbf{c}}_i) \Rightarrow \hat{c}_i) \wedge (\hat{c}_i \Rightarrow \phi_j(\hat{\mathbf{c}}_i)) = \\ &= (\neg\phi_j(\hat{\mathbf{c}}_i) \vee \hat{c}_i) \wedge (\neg\hat{c}_i \vee \phi_j(\hat{\mathbf{c}}_i)) = \\ &= \min\{\max\{1 - \phi_j(\hat{\mathbf{c}}_i), \hat{c}_i\}, \max\{1 - \hat{c}_i, \phi_j(\hat{\mathbf{c}}_i)\}\}\end{aligned}$$

$$\hat{y}_j = \min_{i=1}^k \{\max\{1 - \psi_j(\hat{\mathbf{c}}_i), \ell_{ji}\}\}$$



## Making rules parsimonious

$$\gamma_{ji} = \log \left( \frac{\exp(\text{MLP}_j(\hat{\mathbf{c}}_i))}{\sum_{i'=1}^k \exp(\text{MLP}_j(\hat{\mathbf{c}}_{i'}))} \right)$$
$$r_{ji} = \psi_j(\hat{\mathbf{c}}_i) = \sigma \left( \gamma_{ji} - \frac{1}{k} \sum_{i'=1}^k \gamma_{ji'} \right)$$

# DCR outperforms interpretable models




Figure 3. Mean ROC AUC for task predictions for all baselines across all tasks (the higher the better). DCR often outperforms interpretable concept-based models. *CE* stands for concept embeddings, while *CT* for concept truth degrees. Models trained on concept embeddings are not interpretable as concept embeddings lack a clear semantic for individual embedding dimensions.

# DCR matches the accuracy of neural-symbolic systems trained using human rules

*Table 1.* Task accuracy on the *MNIST-addition* dataset. The neural-symbolic baselines use the knowledge of the symbolic task to distantly supervise the image recognition task. DCR achieves similar performances even though it learns the rules from scratch.

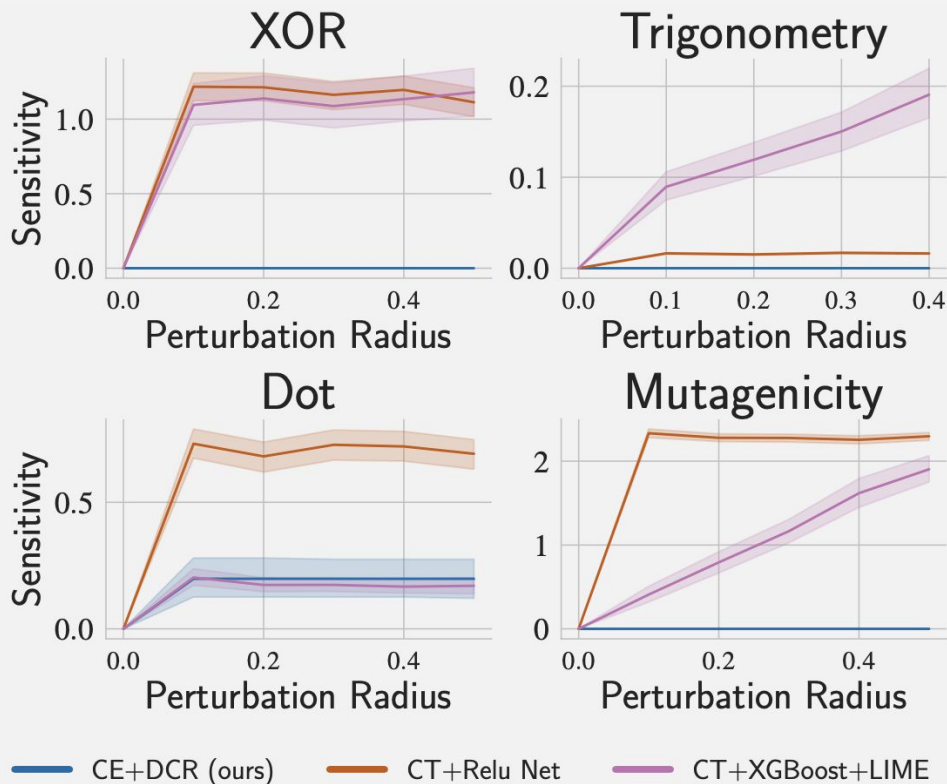
MODEL	ACCURACY (%)
With ground truth rules	
DeepProbLog	$97.2 \pm 0.5$
DeepStochLog	$97.9 \pm 0.1$
Embed2Sym	$97.7 \pm 0.1$
LTN	$98.0 \pm 0.1$
Without ground truth rules	
DCR(ours)	$97.4 \pm 0.2$

# DCR discovers semantically meaningful logic rules

Table 2. Error rate of Booleanised DCR rules w.r.t. ground truth rules. Error rate represents how often the label predicted by a Booleanised rule differs from the fuzzy rule generated by our model. The error rate is reported with the mean and standard error of the mean. A full list of logic rules for MNIST is in Appendix [H](#).

GROUND-TRUTH RULE	PREDICTED RULE	ERROR (%)
<b>XOR</b>		
$y_0 \leftarrow \neg c_0 \wedge \neg c_1$	$y_0 \leftarrow \neg c_0 \wedge \neg c_1$	$0.00 \pm 0.00$
$y_0 \leftarrow c_0 \wedge c_1$	$y_0 \leftarrow c_0 \wedge c_1$	$0.00 \pm 0.00$
$y_1 \leftarrow \neg c_0 \wedge c_1$	$y_1 \leftarrow \neg c_0 \wedge c_1$	$0.02 \pm 0.02$
$y_1 \leftarrow c_0 \wedge \neg c_1$	$y_1 \leftarrow c_0 \wedge \neg c_1$	$0.01 \pm 0.01$
<b>Trigonometry</b>		
$y_0 \leftarrow \neg c_0 \wedge \neg c_1 \wedge \neg c_2$	$y_0 \leftarrow \neg c_0 \wedge \neg c_1 \wedge \neg c_2$	$0.00 \pm 0.00$
$y_1 \leftarrow c_0 \wedge c_1 \wedge c_2$	$y_1 \leftarrow c_0 \wedge c_1 \wedge c_2$	$0.00 \pm 0.00$
<b>MNIST-Addition</b>		
$y_{18} \leftarrow c'_9 \wedge c''_9$	$y_{18} \leftarrow c'_9 \wedge c''_9$	$0.00 \pm 0.00$
$y_{17} \leftarrow c'_9 \wedge c''_8$	$y_{17} \leftarrow c'_9 \wedge c''_8$	$0.00 \pm 0.00$
$y_{17} \leftarrow c'_8 \wedge c''_9$	$y_{17} \leftarrow c'_8 \wedge c''_9$	$0.00 \pm 0.00$

# DCR rules are stable under small perturbations



## DCR explains prediction mistakes

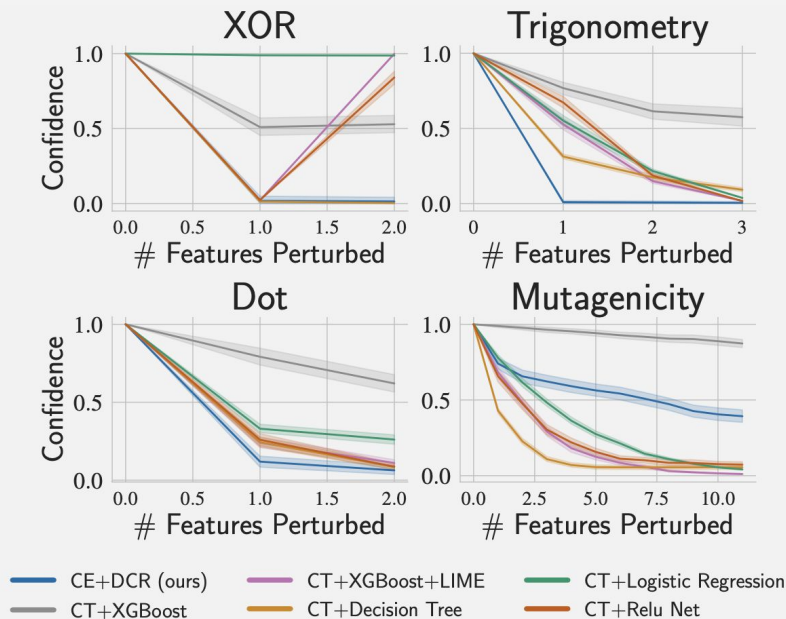
*Table 3.* DCR explains prediction errors.

<b>Dataset</b>	<b>Concepts</b>	<b>DCR rule</b>	<b>Ground truth label</b>
XOR	[0.0, 0.0]	$y = 0 \leftarrow \neg c_0 \wedge \neg c_1$	$y = 1$
Trigonometry	[0.0, 1.0, 1.0]	$y = 1 \leftarrow \neg c_0 \wedge c_1 \wedge c_2$	$y = 0$
Trigonometry	[0.0, 1.0, 0.0]	$y = 0 \leftarrow \neg c_0 \wedge c_1 \wedge \neg c_2$	$y = 1$
Trigonometry	[0.0, 1.0, 1.0]	$y = 1 \leftarrow \neg c_0 \wedge c_1 \wedge c_2$	$y = 0$
Trigonometry	[0.0, 1.0, 1.0]	$y = 1 \leftarrow \neg c_0 \wedge c_1 \wedge c_2$	$y = 0$

# Counterfactual Examples

- generate counter-examples **as close as possible to the original sample**  $|x - x^\star| < \epsilon$
- first **rank** the concepts present in the rule according to their **relevance scores**
-

# DCR enables discovering counterfactual examples



*Figure 5.* Model confidence as a function of the number of perturbed features on counterfactual examples. The lower, the better. Similarly to interpretable methods, DCR prediction confidence quickly drops after inverting the truth degree of a small set of relevant concepts, facilitating the discovery of counterfactual examples.



# Neuro-Symbolic Architecture

This method was based on **Neural**<sub>Symbolic</sub>

This uses a neural net that is generated from symbolic rules. An example is the **Neural Theorem Prover**, which constructs a neural network from an AND-OR proof tree generated from knowledge base rules and terms. **Logic Tensor Networks** also fall into this category [Wikipedia].

# Let's Discuss Limitations of this work and how to improve it

Brainstorm