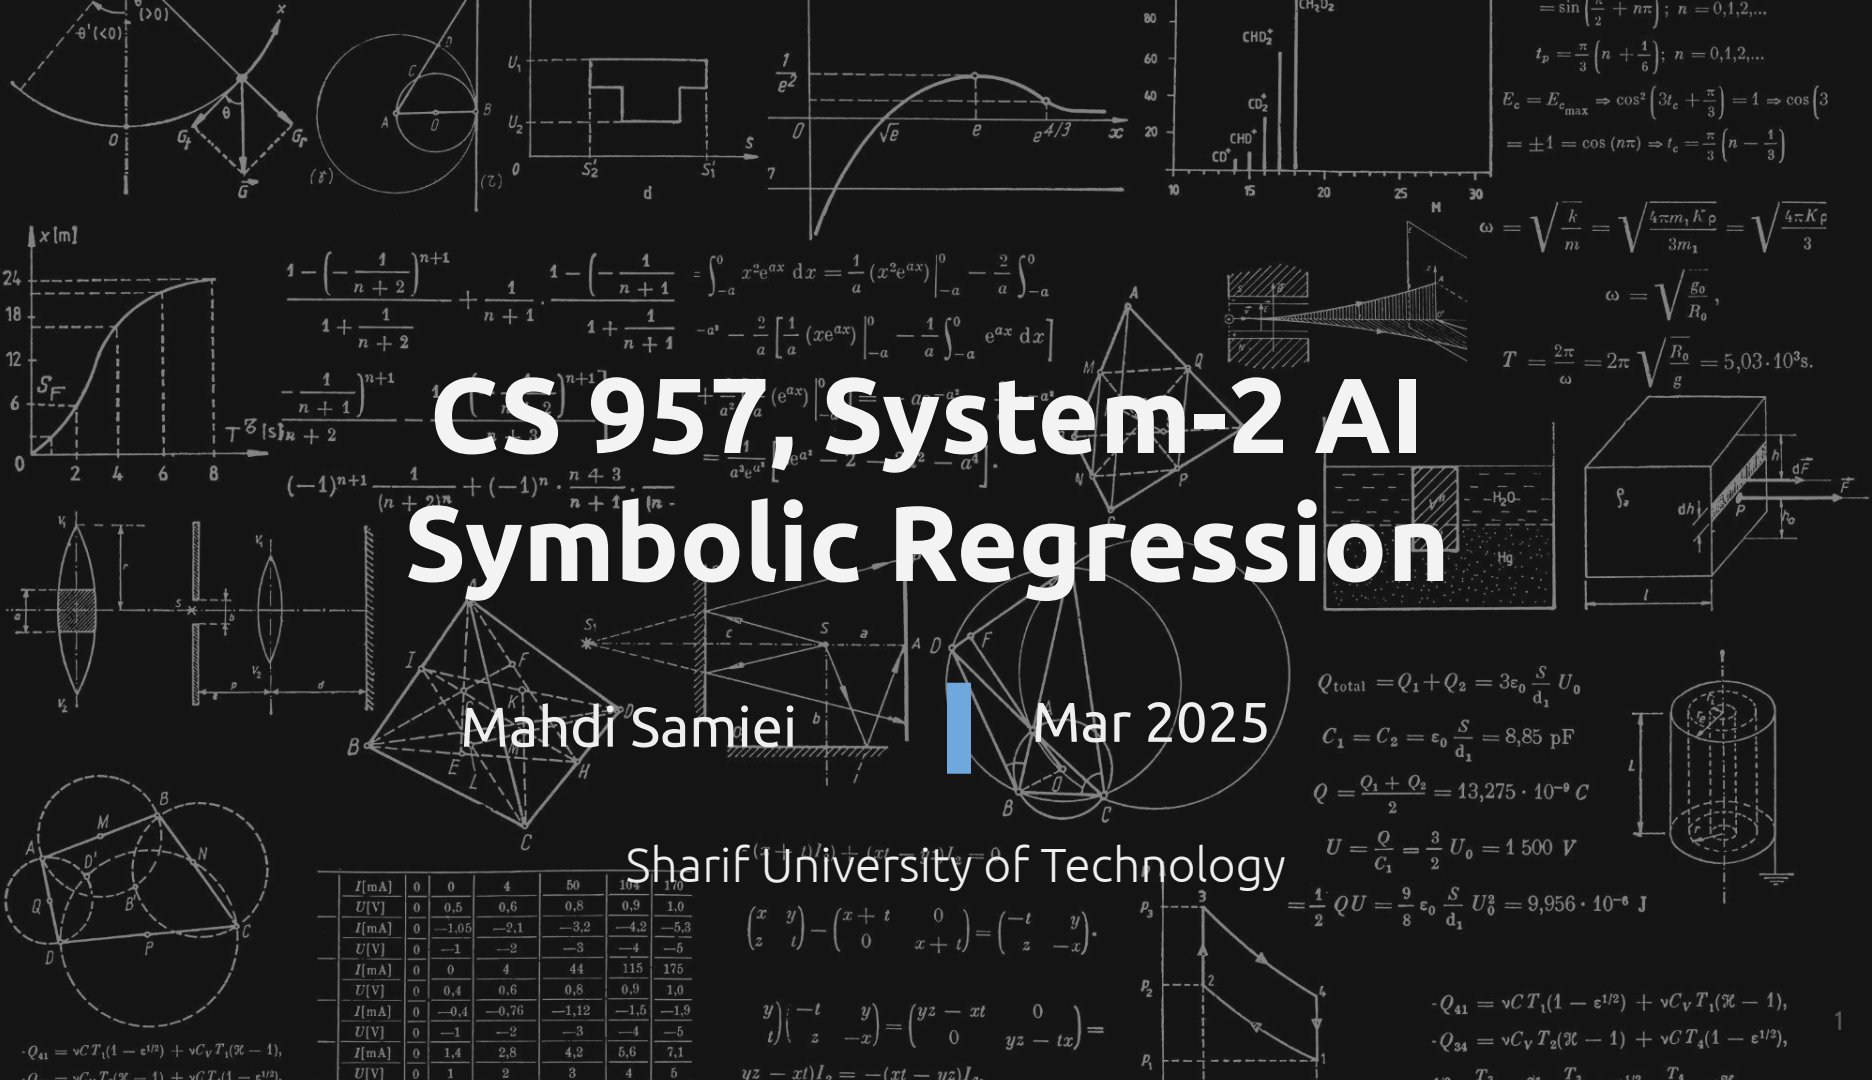


# CS 957, System-2 AI Symbolic Regression

Mahdi Samiei

Mar 2025

Sharif University of Technology




$$-Q_{41} = \sqrt{C} T_1 (1 - \epsilon^{1/2}) + \sqrt{C_V} T_1 (\mathcal{H} - 1),$$

$I$ [mA]	0	0	4	50	100	150
$U$ [V]	0	0.5	0.6	0.8	0.9	1.0
$I$ [mA]	0	-1.05	-2.1	-3.2	-4.2	-5.3
$U$ [V]	0	-1	-2	-3	-4	-5
$I$ [mA]	0	0	4	44	115	175
$U$ [V]	0	0.4	0.6	0.8	0.9	1.0
$I$ [mA]	0	-0.4	-0.76	-1.12	-1.5	-1.9
$U$ [V]	0	-1	-2	-3	-4	-5
$I$ [mA]	0	1.4	2.8	4.2	5.6	7.1
$U$ [V]	0	1	2	3	4	5

$$\begin{pmatrix} x & y \\ z & t \end{pmatrix} - \begin{pmatrix} x+t & 0 \\ 0 & x+t \end{pmatrix} = \begin{pmatrix} -t & y \\ z & -x \end{pmatrix}.$$

$$y \begin{pmatrix} -t & y \\ z & -x \end{pmatrix} = \begin{pmatrix} yz - xt & 0 \\ 0 & yz - tx \end{pmatrix} =$$

$$yz - xt) I_2 = -(xt - yz) I_2.$$



$$Q_{\text{total}} = Q_1 + Q_2 = 3\epsilon_0 \frac{S}{d_1} U_0$$

$$C_1 = C_2 = \epsilon_0 \frac{S}{d_1} = 8,85 \text{ pF}$$

$$Q = \frac{Q_1 + Q_2}{2} = 13,275 \cdot 10^{-9} \text{ C}$$

$$U = \frac{Q}{C_1} = \frac{3}{2} U_0 = 1500 \text{ V}$$

$$= \frac{1}{2} Q U = \frac{9}{8} \epsilon_0 \frac{S}{d_1} U_0^2 = 9,956 \cdot 10^{-6} \text{ J}$$

$$-Q_{41} = \sqrt{C} T_1 (1 - \epsilon^{1/2}) + \sqrt{C_V} T_1 (\mathcal{H} - 1),$$

$$-Q_{34} = \sqrt{C_V} T_2 (\mathcal{H} - 1) + \sqrt{C} T_3 (1 - \epsilon^{1/2}),$$


# A Simple Question ...

If you were given a list of numbers like this, what would you predict?


X	Y
0.31	0.0961
-0.1	0.01
0.21	0.0441
0.64	0.4096
-0.32	0.1024
0.9	0.81
0.5	0.25
0.2	?
1.1	?

# Let's Fit a Neural Model ...


## Seems Good!



# Can't Extrapolate :(




Can't Extrapolate :(



# More Training Data? More Parameters?

X	Y
0.31	0.0961
1.1	1.21
0.21	0.0441
5	25
-0.32	0.1024
9.1	82.81
6.3	39.69
1.2	?
10.1	?

But there is always another ood :(((



# What is Symbolic Regression



discover a mathematical expression that describes the relationship between input output variables.




not a simple task!



Scientific discovery, Engineering system identification, Time series prediction, ...




Remember NP? Hard to Answer, Easy to Verify





# What is Symbolic Regression, vs Neural Regression

 Traditional regression fits parameters to a predefined model structure, while symbolic regression searches for both the structure and parameters simultaneously.

 Interpretable formula

 Extrapolation to Out-of-Distribution Data

 Generalization to more complex formulas

 Regression is a metaphor, get the gist

# What is Symbolic Regression, Structure of Task

➡ Consider Inference Phase: Given  $\{(x_i, y_i)\}$ , We should discover  $f$  that  $y_i = f(x_i)$

➡ Indeed each task is a sample for symbolic regression.

➡ Similar to *Meta-Learning* Setting

🤔 Remember Definition of Generalization?

📊 The performance of our symbolic regression method can be measured either on OoD data or OoD equations.

# What is Symbolic Regression, Structure of Data


🌳 We usually represent the equation in the form of a computational tree.

🌳 Symbolic Regression Two Phases:

① Discover Skeleton with Placeholders for Constants

② Find Best Value for Constants by Optimization

😜 What is Most Dummy Idea?



# Dummy Idea: Brute Force, Search


 Brute-Force, Exhaustive Search

 Local Search Methods, User Heuristic Function

-  Definition of Successor Function

-  Discrete Nature

-  Local Minimas





# Better Idea, Genetic Algorithm



Unrivaled until 2016



Effective even today




The reason?

Combining sub-solutions.



Problems?



## Main drawbacks of classic methods

- ✗ string of symbols grows exponentially with the string length.
- ✗ They do not improve with experience. As every equation is regressed from scratch, the system does not improve if access to more data from different equations is given.
- ✗ The inductive bias is opaque. It is difficult for the user to steer the prior towards a specific class of equations (e.g. polynomials, etc.).

# First Learning Based Ideas:

They remind me of Neurosymbolic!


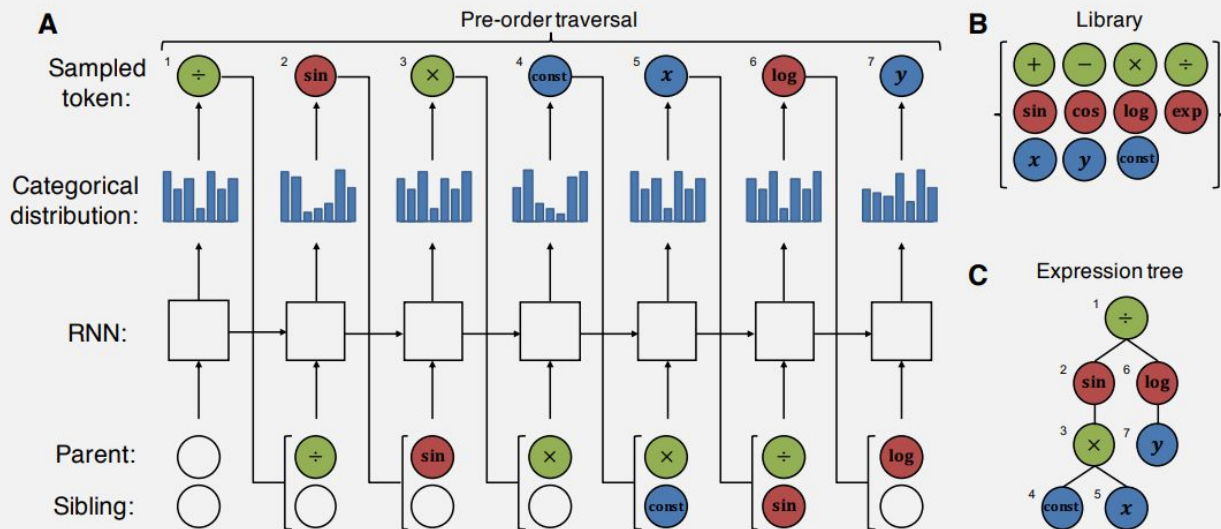


Figure 1: Network architecture of the proposed Equation Learner (EQL) for 3 layers ( $L = 3$ ) and one neuron per type ( $u = 4, v = 1$ ).

Martius, G. and Lampert, C. H. Extrapolation and learning Equations. 2016

# More Advanced Ideas, A RNN for a Task

the network has to be retrained from scratch for each new equation



Petersen, B. K. Deep symbolic regression: Recovering mathematical expressions from data via risk-seeking policy gradients. International Conference on Learning Representations, 2021



# Neural Symbolic Regression that Scales

Luca Biggio, Tommaso Bendinelli, Alexander Neitz,  
Aurelien Lucchi, Giambattista Parascandolo  
June 2021

# Main drawbacks of classic methods

✗ string of symbols grows exponentially with the string length.

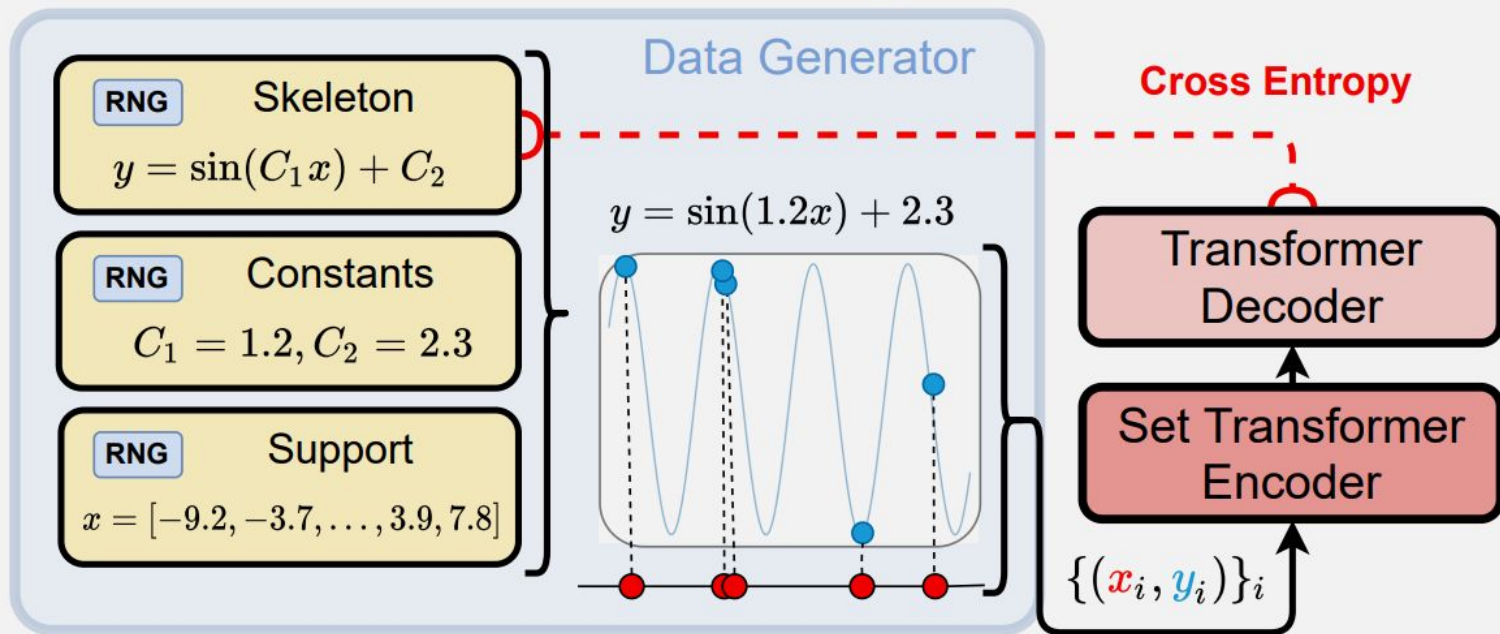
✗ They do not improve with experience. As every equation is regressed from scratch, the system does not improve if access to more data from different equations is given.

✗ The inductive bias is opaque. It is difficult for the user to steer the prior towards a specific class of equations (e.g. polynomials, etc.).

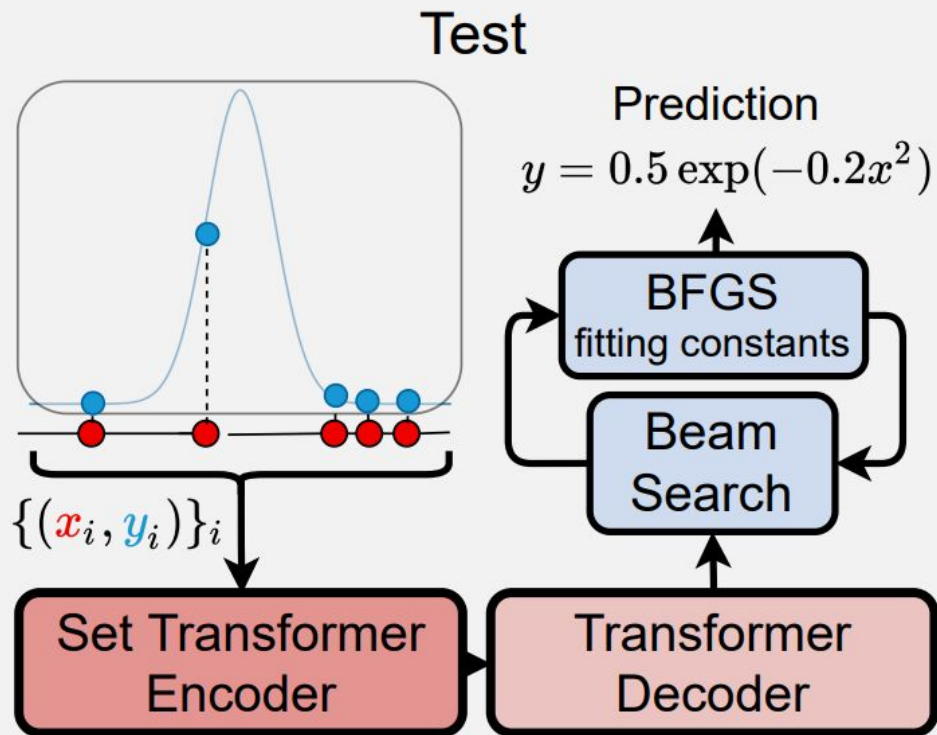
✓ **we show that a strong symbolic regression can be purely learned from data**

# Pre-Training


## Pre-Training



# Test Time



# Evaluation, Scale Train Time



# Evaluation, Scale Test Time

Table 1. Hyper-parameters that vary to increase the amount of compute invested by every method.

Method	Hyper-param	Range
G. Proc. (Rasmussen 2003)	Opt. restarts	$\{8, 16, 32\}$
Genetic Prog. (Koza 1994)	Pop. size	$\{2^{10}, \dots, 2^{17}\}$
DSR (Petersen 2021)	Epochs	$\{2^2, \dots, 2^7\}$
NeSymReS (ours)	Beam size	$\{2^0, \dots, 2^8\}$

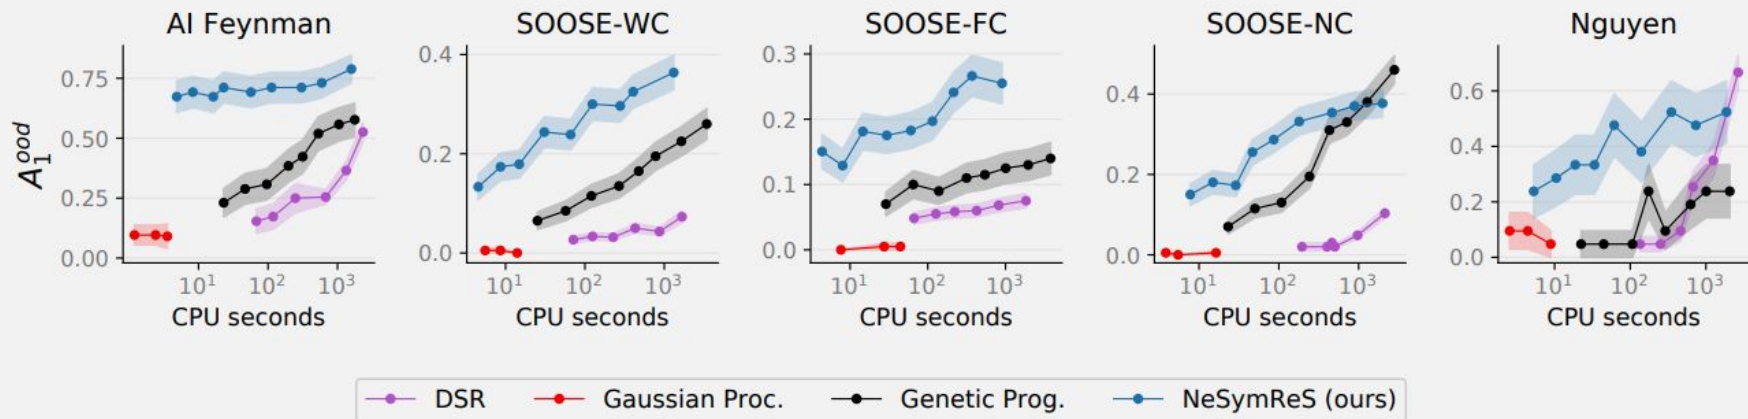
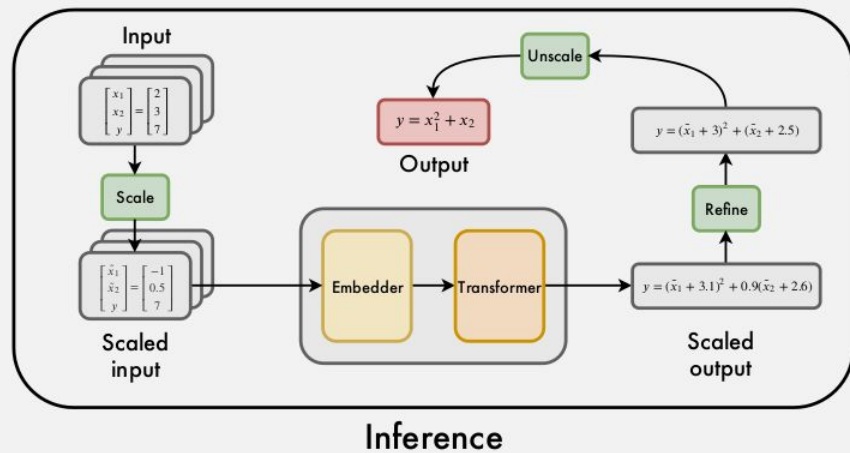
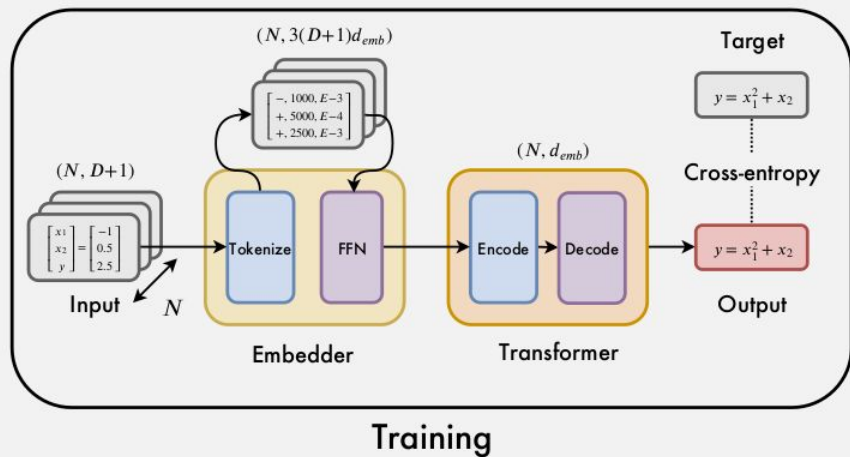


Figure 4. Accuracy out of distribution as a function of time for all methods ran on a single CPU per equation.

# End-to-End Symbolic Regression with Transformers

Pierre-Alexandre Kamienny, Stéphane d'Ascoli, Guillaume Lample, François Charton  
April 2022

# Learn to Predict End-to-End





# Inference tricks

Model	Function $f(x, y)$
Target	$\sin(10x) \exp(0.1y)$
Skeleton + BFGS	$-\sin(1.7x)(0.059y + 0.19)$
E2E no BFGS	$\sin(9.9x) \exp(0.1y)$
E2E + BFGS random init	$-\sin(0.095x) \exp(0.27y)$
E2E + BFGS model init	$\sin(10x) \exp(0.1y)$

**Table 1: The importance of an end-to-end model with refinement.** The skeleton approach recovers an incorrect skeleton. The E2E approach predicts the right skeleton. Refinement worsens original prediction when randomly initialized, and yields the correct result when initialized with predicted constants.

# LLM-SR: Scientific Equation Discovery via Programming with Large Language Models

Parshin Shojaee, Kazem Meidani, Shashank Gupta, Amir Barati Farimani, Chandan K Reddy  
April 2024

# Get Rid of Computation Tree

all  $i$ . The discovered equation should not only accurately fit the observed data points but also exhibit strong generalization capabilities to unseen data while maintaining interpretability.

Current SR methods typically represent equations using techniques such as expression trees (Cranmer, 2023), prefix sequences (Petersen et al., 2021; Biggio et al., 2021), or context-free grammars (Brence et al., 2021), constructing a search space  $\mathcal{F}$ . These representations provide limited and structured search spaces, enabling evolutionary search approaches like genetic programming (GP) to explore and find candidate expressions. In contrast, our work employs program functions to directly map inputs  $x$  to output targets  $y$ : `def f(x): ... return y`.

# LLM-SR

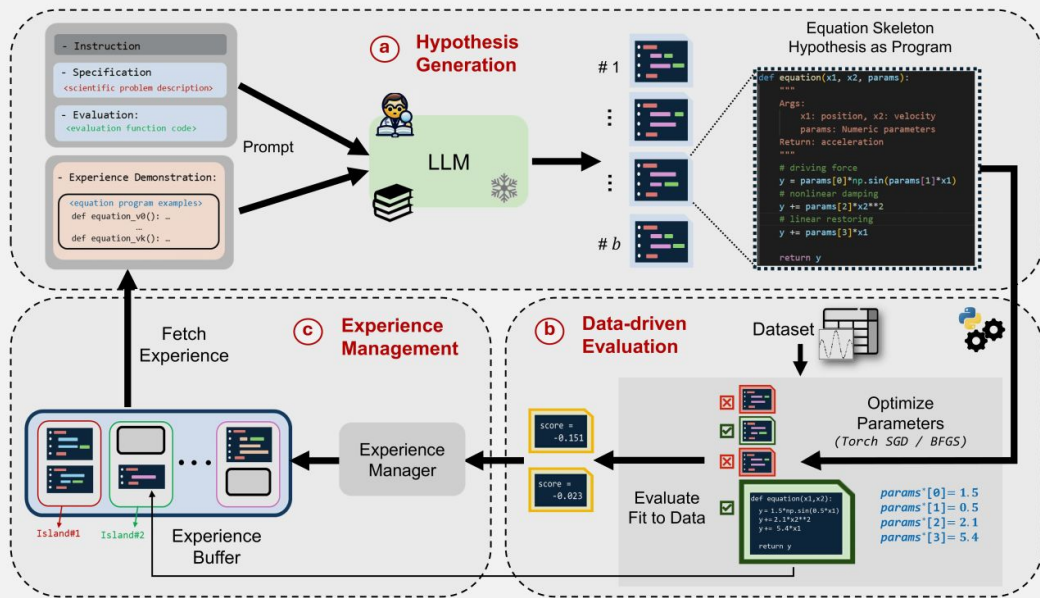


Figure 1: **The LLM-SR framework**, consisting of three main steps: **(a) Hypothesis Generation**, where LLM generates equation program skeletons based on a structured prompt; **(b) Data-driven Evaluation**, which optimizes the parameters of each equation skeleton hypothesis and assesses its fit to the data; and **(c) Experience Management**, which maintains a diverse buffer of high-scoring hypotheses to provide informative in-context examples into LLM's prompt for iterative refinement.

# LLM-SR, Prompt

Can you give me the complete version of last function?

Instruction

Consider physical meaning and relationships of inputs while completing the function.

"""

*Find the mathematical function skeleton that represents E. Coli bacterial growth  
→ rate, given data on the population density of bacterial species, substrate  
→ concentration, temperature, and pH level.*

"""

Problem  
Specification

```
import numpy as np
from scipy.optimize import minimize
```

```
def loss_function(params, x0, x1, x2, x3, y_true):
    y_pred = equation(x0, x1, x2, x3, params)
    return np.mean((y_pred - y_true) ** 2)
```

Evaluation and  
Optimization

```
def evaluate(data: dict) -> float:
    """Evaluate equation on input and output observations."""
    # Load true data observations
    inputs, outputs = data['inputs'], data['outputs']
    # Optimize equation skeleton parameters
    loss_partial = lambda params: loss_function(params, *inputs.T[:4], outputs)
    params_initial_guess = [1.0]*P
    result = minimize(loss_partial, params_initial_guess, method='BFGS')
    optimized_params = result.x
    # Return evaluation score
    score = loss_function(optimized_params, *inputs.T[:4])
    return -score if not np.isnan(score) and not np.isinf(score) else None
```

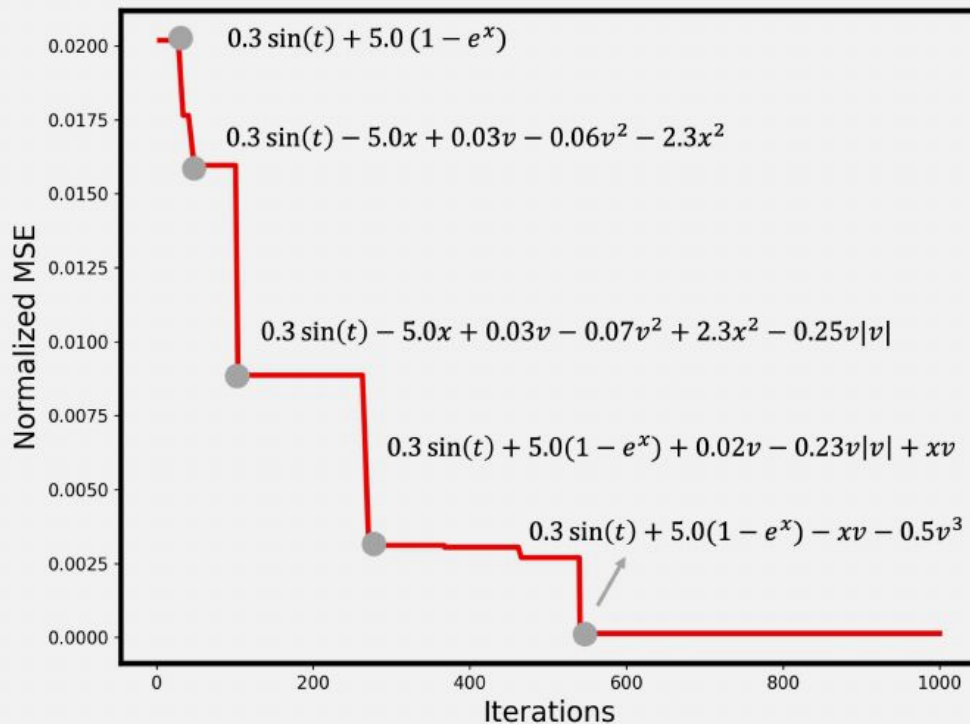
```
def equation_v0(x0: np.ndarray, x1: np.ndarray, x2: np.ndarray,
x3: np.ndarray, params: np.ndarray) -> np.ndarray:
    """Find mathematical function skeleton for the bacterial growth rate.
    Args:
        x0: Population density of the bacterial species.
        x1: Substrate concentration.
        x2: Temperature.
        x3: pH level.
        params: Array of numeric parameters to be optimized
    Return:
        A numpy array representing bacterial growth rate
    """
    return params[0]*x0 + params[1]*x1 + params[2]*x2 + params[3]*x3 + params[4]
```

Equation Program  
Example

```
def equation_v1(x0: np.ndarray, x1: np.ndarray, x2: np.ndarray,
x3: np.ndarray, params: np.ndarray) -> np.ndarray:
    """Improved version of equation_v0"""
```

Function to Complete

# LLM-Guided Search



# LLM-SR Discovers more Accurate Equations

Model	Oscillation 1		Oscillation 2		E. coli growth		Stress-Strain		All	
	ID↓	OOD↓	ID↓	OOD↓	ID↓	OOD↓	ID↓	OOD↓	ID↓	OOD↓
GPlearn	0.0155	0.5567	0.7551	3.188	1.081	1.039	0.1063	0.4091	0.4894	1.298
DSR (Petersen et al., 2021)	0.0087	0.2454	0.0580	0.1945	0.9451	2.4291	0.3326	1.108	0.3361	0.9942
uDSR (Landajuela et al., 2022)	0.0003	0.0007	0.0032	0.0015	0.3322	5.4584	0.0502	0.1761	0.0964	1.409
PySR (Cranmer, 2023)	0.0009	0.3106	0.0002	0.0098	0.0376	1.0141	0.0331	0.1304	0.0179	0.3662
LLM-SR (Mixtral)	<b>7.89e-8</b>	<b>0.0002</b>	0.0030	0.0291	<b>0.0026</b>	<b>0.0037</b>	<b>0.0162</b>	0.0946	<b>0.0054</b>	0.0319
LLM-SR (GPT-3.5)	4.65e-7	0.0005	<b>2.12e-7</b>	<b>3.81e-5</b>	0.0214	0.0264	0.0210	<b>0.0516</b>	0.0106	<b>0.0196</b>

Table 1: **Quantitative performance comparison** of LLM-SR (with GPT-3.5 and Mixtral backbones), and SR baseline models on different scientific benchmark problems measured by Normalized Mean Squared Error. 'All' column indicates average scores over all datasets. SR baselines ran for over 2M iterations, while LLM-SR variants ran for around 2K iterations.



# LLM-SR has better OOD generalization




Figure 3: Comparison of *E. coli* growth rate distributions from LLM-SR, PySR, and uDSR. LLM-SR aligns well with the ground truth, even for OOD data (unshaded regions), demonstrating better generalization than PySR and uDSR, which overfit the ID data (shaded regions and black points).



# LLM-SR Discovers Equations More Efficiently

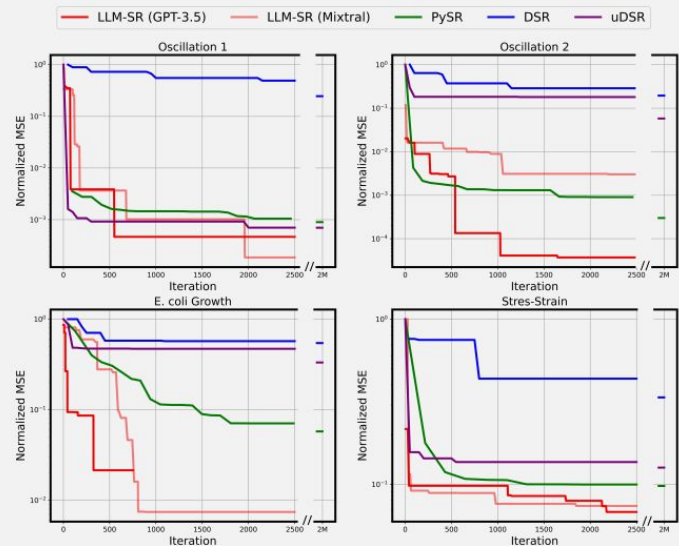
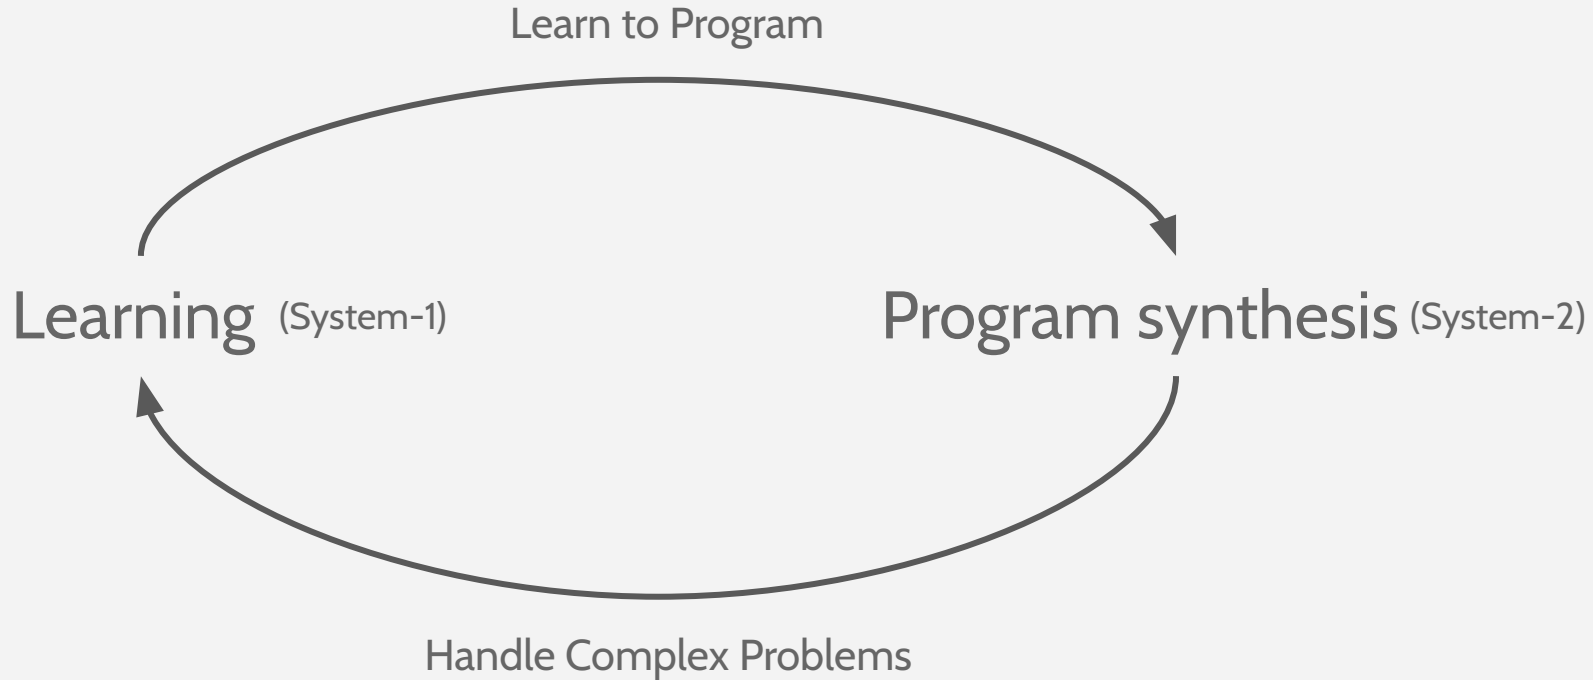


Figure 4: **Best score trajectories** of LLM-SR with GPT-3.5 and Mixtral against SR baselines across different benchmark problems. LLM-SR discovers accurate equations more efficiently, requiring fewer iterations to outperform baselines, which fail to match LLM-SR even after 2M iterations.

# Symbolic Regression: Equation as a Program



# Next Session:

## Program Synthesis, beyond Regression